

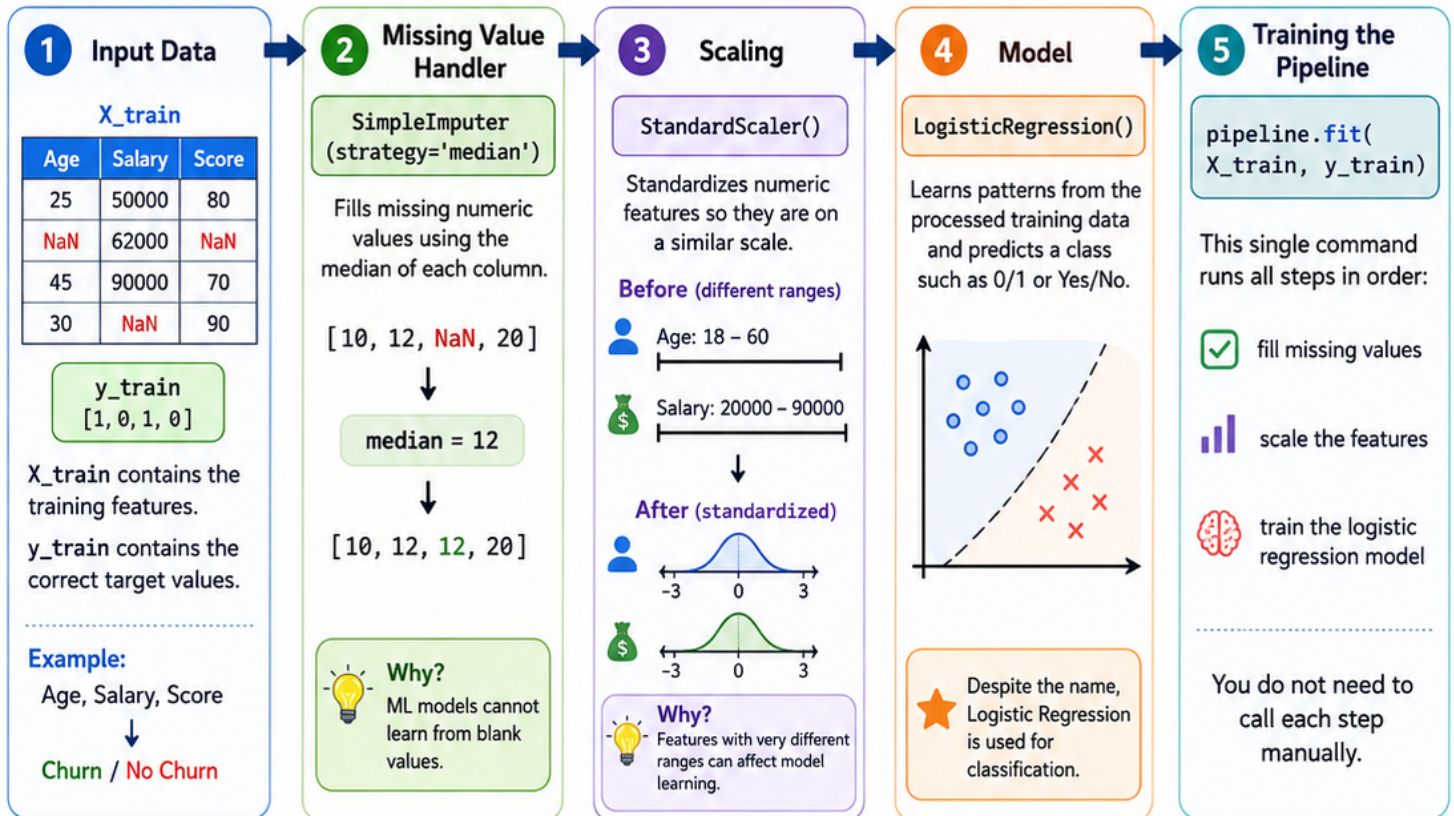
How a scikit-learn Pipeline Works

A beginner-friendly explanation of this code

```
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression

pipeline = Pipeline(steps=[
    ("missing_value_handler", SimpleImputer(strategy="median")),
    ("scaler", StandardScaler()),
    ("model", LogisticRegression())
])

pipeline.fit(X_train, y_train)
```



Why use a Pipeline?



Reusable:
same workflow can be run again.



Cleaner code:
all steps in one place.



Consistent:
same preprocessing every time.



Less error-prone:
fewer manual mistakes.



Beginner Tip: This simple pipeline works best when all input features are numeric. If your dataset also has text or categorical columns, you usually combine pipelines with a **ColumnTransformer**.